

# Concurrency Patterns

(Mostly) for Library Authors

26 May 2015

Aaron Schlesinger  
Sr. Engineer, Iron.io

# About me

- Database & backend systems engineer at Iron.io
- Writing Go for ~2 yrs, JVM for a while during/before that
- Distributed systems at Zynga, StackMob, PayPal, now Iron
- C -> C++ -> Python/PHP -> -> Java -> Scala -> Go. I'm finally happy

# "It's almost too easy to write concurrency bugs in Go"

Because Go has powerful concurrency primitives.

I'm discussing today how to use them responsibly.

I said library authors, but this talk is aimed toward anyone who writes reusable code.

# Today

Conventions & patterns for:

- mutex, chan, select, and `sync.WaitGroup`
- [Timeouts](http://blog.golang.org/go-concurrency-patterns-timing-out-and) (<http://blog.golang.org/go-concurrency-patterns-timing-out-and>), cancellation and [net.Context](http://godoc.org/golang.org/x/net/context) (<http://godoc.org/golang.org/x/net/context>)
- For-select loops

# Something old

Locks have a time and a place. We found one at Iron.io (building a distributed DB).

If you use locks:



But prefer sharing memory by communicating.

## Go-specific conventions

- defer unlocking wherever possible
- Document ownership in your **public** interface
- If mutual exclusion crosses a public interface boundary, abstract it

Also, good conventions already exist from other communities (e.g. lock order conventions).

# Documentation

```
package main

import "errors"

var ErrNotReserved = errors.New("not reserved")

// RequestState is an in-memory table to track the state of all requests.
// States can be any positive number.
type RequestState interface {
    // GetAndReserve waits for the state to be available, then returns it.
    // After this call returns, you have ownership of this state. Use
    // the returned string as an ownership token to indicate your ownership of id
    // in later calls. Returns an error if communication with underlying storage fails.
    // If this func returns error, you do not have ownership.
    GetAndReserve(id string) (int, string, error)

    // SetReserved sets the state of the given request ID, given the ownership ID.
    // doesn't set the state and returned ErrNotReserved if the state wasn't reserved
    // by the given ownership ID. After this call returns successfully, you no
    // longer have ownership of id and ownershipID is invalidated.
    SetReserved(id, ownershipID string, state int) error
}
```

# Channels

Share memory by communicating.

This means passing messages.

- Channels + goroutines are "easy" but powerful enough to build real systems
- They're first class for a reason. Use them by default
- When in doubt, ask why you *shouldn't* use them to communicate between goroutines



# Specifically

- Document channel communication across func boundaries
- Enlist the compiler. Use directional channels
- Don't return a chan unless the func is a **generator** (<https://talks.golang.org/2012/concurrency.slide#25>)
- `close` is a useful signal to callers. Use it and document it

# Example

```
package main

import "sync"

// WatchChanges will watch the state of the given request. ch will send after
// each request state change and will be closed after the request is removed from
// the request state database.
//
// WatchChanges sends on ch from the same goroutine as the caller and will deadlock
// if there is no goroutine already listening on it.
//
// Returns ErrNotFound if the request is not reserved at call time.
// If any error is returned, WatchChanges will neither do any operations on ch.
func WatchChanges(reqID string, ch chan<- int) error

// WatchAll watches for all events on the given request.
//
// The WaitGroup will be done after the request is reserved, and the channel
// will send on each state change, then be closed when the request is released.
//
// The channel will send from a new, internal goroutine, which you are not responsible
// for managing.
func WatchAll(reqID string) (*sync.WaitGroup, <-chan int)
```

# Interlude 1

- Documentation may establish contracts or invariants that code can't or won't
- That being said, the remainder of this talk shows has mostly runnable code
- Imagine that it has proper documentation in it

# Select

## Channels of channels:

- We have (probably more than) a few posts on how/why to use them
- Particularly, listening on channels-of-channels in a `select` enables a goroutine to offer many "services"
- This pattern can be helpful for introspecting the state of a goroutine

# Introspection

Trivial example:

```
func doWork(calc <-chan chan<- int, cur <-chan chan<- int, history <-chan chan<- []int, stats chan<- int) {  
    var all []int  
    i := 0  
    for {  
        select {  
        case ch := <-calc:  
            i := rand.Int()  
            ch <- i  
            all = append(all, i)  
        case ch := <-cur:  
            ch <- i  
        case ch := <-history:  
            ch <- all  
        }  
        i++  
        stats <- i  
    }  
}
```

[Run](#)

# WaitGroup

In general, if you're waiting for a chan to close or receive a `struct {}`, think about using a `WaitGroup`.

Passing these around as signals of behavior can also help you write deterministic tests.

# Notification of an event

Trivial example.

```
package main

import "sync"

// startLoop starts a loop in a goroutine. the returned WaitGroup is done
// after the first loop iteration has started
func startLoop(n int) *sync.WaitGroup {
    var wg sync.WaitGroup
    go func() {
        first := true
        for {
            if first {
                wg.Done()
                first = false
            }
            // do some work here
        }
    }()
    return &wg
}
```

# Revisiting fan-in



(A better image is at <https://talks.golang.org/2012/concurrency.slide#28>

(<https://talks.golang.org/2012/concurrency.slide#28>)



# Why?

If you can do work concurrently, do it. There's no excuse not to.

How I'm defining fan-in:

- A pattern to gather results from 2 or more goroutines
- A procedure you can follow to take sequential code and convert it to concurrent

# Details

- Read about it under "Fan-out, fan-in" section at <https://blog.golang.org/pipelines> (<https://blog.golang.org/pipelines>).
- `sync.WaitGroup` and a few channels make fan-in simple & understandable enough to convert non-concurrent code to concurrent code with this pattern.
- In many cases, you can get an easy latency win without changing an API.

# Sequential datastore queries

```
func GetAll() []int {  
    ints := make([]int, 10)  
    for i := 0; i < 10; i++ {  
        ints[i] = datastoreGet()  
    }  
    return ints  
}
```

[Run](#)

# Concurrent datastore queries

```
func GetAll() []int {  
    wg, ch := getWGAndChan(10) // get a waitgroup that has 10 added to it, and a chan int  
    for i := 0; i < 10; i++ {  
        c := make(chan int)  
        go datastoreGet(c) // sends an int on c then closes after sleeping <= 1 sec  
        go func() {  
            defer wg.Done() // mark this iteration done after receiving on c  
            ch <- <-c        // enhancement: range of c if >1 results  
        }()  
    }  
    go func() {  
        wg.Wait() // wait for all datastoreGets to finish  
        close(ch) // then close the main channel  
    }()  
    ints := make([]int, 10)  
    i := 0  
    for res := range ch { // stream all results from each datastoreGet into the slice  
        ints[i] = res // GetAll can be a generator if you're willing to change API.  
        i++           // that lets you push results back to the caller.  
    }  
    return ints  
}
```

[Run](#)

## Interlude 2

In production, don't grind to a halt when a channel doesn't send.

In Go, we can use `time.Timer` to time out channel receives.

But we're gonna go a step further.

# Timeouts

But, `net.Context` (<https://blog.golang.org/context>) has a nice interface in front of timers.

- I'm cheating here. I said everything would be done with the standard library
- We're *mostly* using contexts as a timer, but it provides a nice interface with a few extra bells & whistles.
- That's my excuse

## Using contexts

- If you are waiting for a channel, don't wait forever
- Ideally, all goroutines take a context
- Contexts add more control and testability

`net.Context` is a good universal tool for timeouts/cancellation in a large codebase.

# Polling a queue

```
func main() {  
    ctx, cancel := context.WithTimeout(context.Background(), 500*time.Millisecond)  
    defer cancel()  
    ch := make(chan string)  
    go poll(ctx, ch) // sends dequeued results on ch. stops & closes ch when ctx.Done() receives  
    for polled := range ch {  
        fmt.Println(polled)  
    }  
}
```

[Run](#)



# Contexts in a distributed system

## The Tail at Scale.

- Jeff Dean talk/paper. I originally saw it at a [Ricon 2013 Talk](https://www.youtube.com/watch?v=C_PxVdQmfpk)
- Hedged requests: do a few identical GET (e.g. no side effects) requests, cancel remaining requests after first returns

Rob showed a variant in <https://talks.golang.org/2012/concurrency.slide#50>

(<https://talks.golang.org/2012/concurrency.slide#50>)

# Adding cancellation

## With net.Context!

```
func main() {  
    ch := make(chan int)  
    ctx, cancel := context.WithTimeout(context.Background(), 10*time.Millisecond)  
    defer cancel()  
    for i := 0; i < 10; i++ {  
        // get (not shown) sleeps for a random duration <= 100ms,  
        // then sends a random int on ch. stops if ctx.Done() receives.  
        go get(ctx, ch)  
    }  
    select {  
    case i := <-ch:  
        fmt.Printf("got result %d\n", i)  
    case <-ctx.Done():  
        fmt.Println("got no result")  
    }  
}
```

[Run](#)

# That was the naïve implementation

But, it's not too hard to get "fancy"

- Don't send 2nd request until 1st is past watermark (95th percentile expected latency)
- Cancel in-flight requests (pass context to RPC subsystem)
- Target-target communication (pass info on other in-flight requests over RPC)

# For-select loops

Running a (possibly infinite) loop inside a goroutine, selecting on 1 or more channels at each iteration.

Lots of possible applications:

- Event loops
- GC
- Synchronizing on state (like an actor)

# Patterns

- ack before and after real work is done. testing is easier and rate limiting/backpressure is easy
- if you're ticking, build acks into `time.Ticker`
- `net.Context` for cancellation
- `sync.WaitGroup` for started and stopped loops

# A for-select poller

```
func poll(ctx context.Context) (*sync.WaitGroup, *sync.WaitGroup, <-chan string) {
    var start, end sync.WaitGroup
    start.Add(1)
    end.Add(1)
    ch := make(chan string)
    go func() {
        defer close(ch)
        defer end.Done()
        start.Done()
        for {
            time.Sleep(5 * time.Millisecond)
            select {
            case <-ctx.Done():
                return
            case ch <- "element " + strconv.Itoa(rand.Int()):
            }
        }
    }()
    return &start, &end, ch
}
```

# Driving the poller

```
func main() {  
    ctx, cancel := context.WithTimeout(context.Background(), 10*time.Millisecond)  
    defer cancel()  
    mainCh, wg := makeThings(10) // make a chan string and a wg that has 10 added to it  
    for i := 0; i < 10; i++ {  
        start, _, ch := poll(ctx)  
        start.Wait()  
        go func() {  
            defer wg.Done()  
            for str := range ch {  
                mainCh <- str  
            }  
        }()  
    }  
    go func() {  
        wg.Wait()  
        close(mainCh)  
    }()  
    printCh(mainCh) // loops on mainCh until it's closed  
}
```

[Run](#)

# Notes

- The poller is missing the ack
- We have a small utility at Iron to add acks to `time.Ticker` and `time.Timer`
- Exercise left to the reader



# Conclusion

Go has really good concurrency abstractions built into the language.

I believe we (the community) are starting to build good patterns & tools to *responsibly* build on them.

# If you take one thing away

Use `net.Context` somewhere in your codebase.

Or, at least try it.

It's simple and powerful, and follows the "Go way."

# If you take two things away

Add reading and understanding [Go Concurrency Patterns](https://talks.golang.org/2012/concurrency.slide) (https://talks.golang.org/2012/concurrency.slide) .

# Thank you

Aaron Schlesinger  
Sr. Engineer, Iron.io

[aaron@iron.io](mailto:aaron@iron.io) (<mailto:aaron@iron.io>)

<http://github.com/arschles> (<http://github.com/arschles>)

[@arschles](http://twitter.com/arschles) (<http://twitter.com/arschles>)

